



Figure 1: This is a picture of one of Pavlov's dogs, it has been stuffed and is preserved at The Pavlov Museum. You can see where the saliva tube and container has been implanted. [Picture from http://en.wikipedia.org/wiki/Ivan_Pavlov]

Introduction

These notes are about the classical conditioning and the mesocortical dopaminergic pathway which is believed to be responsible for the brain's *reward system*.

Classical conditioning

Pavlov is the biggest fool I know; any policeman could tell you that much about a dog. - George Bernard Shaw

In Pavlov's famous experiment, conducted at the turn of the nineteenth and twentieth centuries, a bell is rung a short time before a dog is fed; obviously feeding causes salivation in the dog, but the curious thing is that after a while the dog salivates as soon as it hears the bell. This experimental and the conclusion that were drawn from it were hugely controversial at the time, opinions ranged from Shaw's above, claiming that nothing interesting had been measured or concluded. For its proponents, Pavlovian conditioning seemed to promise a new scientific era of psychology and even promised the 'the perfectibility of man' and featured, for example, in Aldous Huxley's dystopian novel *Brave New World* and his utopian novel *The Island*. As for Shaw:

If 'A' is drowning on one side of a pier and 'B' is equally drowning on the other, and you have one lifebelt, to which of the two would you like to throw it? Which would I save, Pavloff or Shaw? What is the good of Shaw? And what is the good of Pavloff? Pavloff is a star which lights the world, shining above a vista hitherto unexplored. Why should I hesitate with my lifebelt for one moment? - H.G. Wells

These days we describe the bell as an *conditioned stimulus* (CS), it produces salivation only after training, the food is an *unconditioned stimulus* (US), it always produces response. In other words, the US is the food since it already produces the reaction: salivation, whereas the bell is the CS since it only produces the response after training, that is, after conditioning. Classical, or Pavlovian, conditioning is also distinguished from *instrumental* or *operant* conditioning in which the actions of animal determine the reinforcement; like Pavlovian conditioning,




Figure 2: Changes in w . In each trial a stimulus is presented with probability 0.5; during the first 150 trials the stimulus is accompanied by a reward of $r = 1$, after that by $r = 0$. Plotted is the resulting Rescorla-Wagner changes in w as it learns the reward and as the conditioning is extinguished. The learning rate is $\eta = 0.05$. This is roughly based on a similar figure in [2].

operant conditioning has a controversial history, it is associated with another behaviourist, B.F. Skinner, but away from the complex philosophical and political interpretations, both forms of conditioning are now important neuroscientific tools.

Models of classical conditioning

The widely used Rescorla-Wagner model of classical conditioning works [1] like a perceptron; based on the stimulus the animal anticipates a reward and adjusts its prediction according to its accuracy. Hence, if x is a binary value representing the presence or absence of the stimulus, $r > 0$ is the reward; a negative r would correspond to an aversive event, v is the predicted reward and w the weight used by the animal to predict the reward:

$$v = wx \quad (1)$$

The Rescorla-Wagner rule is then

$$w \rightarrow w + \eta \delta x \quad (2)$$

where $\delta = r - v$ is the error in prediction and η is a learning rate. These dynamics are illustrated in 2.

The Rescorla-Wagner rule generalizes to more than one stimulus-reward pair. Say x_i is the binary value representing the presence or absence of the i th stimulus and r_i is the corresponding reward, then the total reward is

$$r = \sum_i x_i r_i \quad (3)$$

and the predicted reward is

$$v = \sum_i x_i w_i \quad (4)$$

and the learning rule is

$$w_i \rightarrow w_i + \eta \delta x_i \quad (5)$$

where $\delta = r - v$ as before.

One significant victory for this proposal is that it explains blocking. Consider conditioning a reward r on a stimulus s_1 and then changing so that there are two stimuli used to predict r , s_1 and s_2 ; now, when s_2 is shown on its own to the animal it does not anticipate the reward. Thus, for example, Pavlov's dog might be shown a light just before it is fed and will soon salivate when it sees the light; next a light is lit and a bell rung before feeding, now, if the bell is rung on its own, the dog does not salivate; the light has *blocked* the bell. This is an easy consequence of the Rescorla-Wagner rule since the w for the light already gives a correct prediction of the target and so the w for the bell stays at zero as there is no error. Blocking is not a consequence of other models of conditioning proposed at the same time. It has, however, been observed in behavior. [3, 4].

Ventral tegmental area

The ventral tegmental area (VTA) is located immediately beside the substantia nigra (SN) in the midbrain. In Wikipedia it says '[VTA] is important in cognition, motivation, orgasm, drug addiction, intense emotions relating to love, and several psychiatric disorders.' We are interested in it here because it is believed to play an important role in the reward system.

The VTA has a large number of dopaminergic neurons, dopamine is a neuromodulator; the level of dopamine alters the dynamics of synapses: different synapses have different dopamine receptors and so their dynamics might be altered in different ways. Dopaminergic neurons are not common, there are about 400,000 in the human brain, about half of these are in the VTA and half of VTA neurons are dopaminergic. these dopaminergic neurons project to diverse areas in the brain, see Fig. 3, including hippocampus, basal ganglia and the prefrontal cortex. These projections transmit dopamine to these areas. Another peculiarity of VTA is that contains a large number of gap junctions.

The idea here is that δ , the error, is calculated in dopaminergic VTA neurons and that neuromodulation produces the Rescorla-Wagner rule. This is sketched out in Fig. 4. Evidence for this can be seen in a famous experiment [5] in which the activity of dopaminergic cells was recorded in monkeys during conditioning. Before condition the dopaminergic neurons fire at an elevated rate when the reward is received, after conditioning they fire at a depressed rate if the anticipated reward fails to appear after the stimulus, Fig. 5.

Another thing is apparent from Fig. 5: when the animal has been conditioned the dopaminergic neurons fire at the stimulus, the activity shifts forward from reward to the stimulus that predicts it. In a way this makes sense, the unanticipated event is the stimulus; this predicts the reward but is itself a surprise. This is also obviously useful, it allows the credit for the reward to be usefully associated with the event that predicts it. If there are a series of events one predicting the other it allows the credit to filter forwards to the event that initiates the whole sequence, this might have consequence for behavior. A mechanism for this shifting forward is known, it is a sort of time-segmented Rescorla-Wagner called temporal difference learning [6, 7].




Figure 3: This shows the dopamine pathways from VTA and from SN; VTA has major projections to hippocampus, the nucleus accumbens in the basal ganlia and to the prefrontal cortex. [Picture from <http://en.wikipedia.org/wiki/Dopamine>]

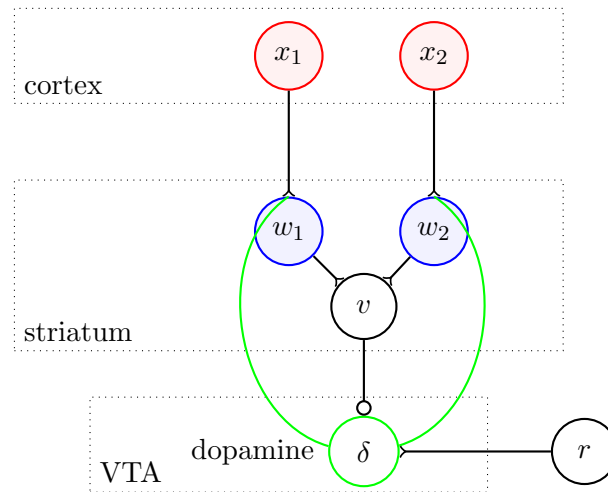


Figure 4: A schematic of the VTA reward circuit. The conditioned stimulus is presented and this is communicated via the cortex, neurons in the striatum adjust this to give the input $w_i x_i$ and these are added producing the estimated reward. This inhibits a dopaminergic neuron in VTA, this neuron also receives excitatory input corresponding to the actual reward, this difference is δ and δ effects dopamine modulation w_1 and w_2 effecting some version of the Rescorla-Wagner rule.

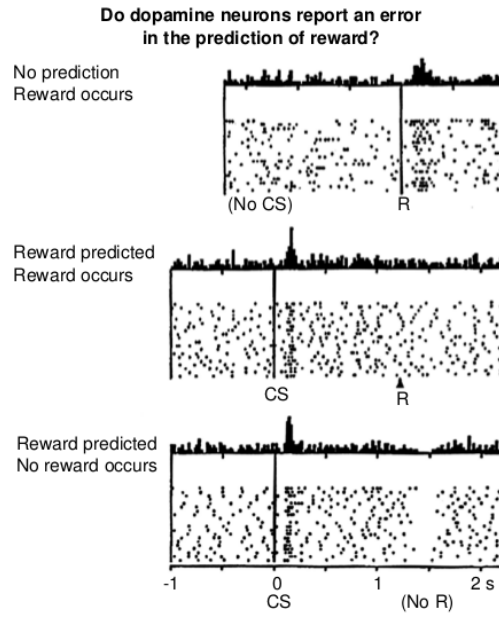


Figure 5: Dopaminergic cell activity. In the top panel there has been no conditioning, this means the predicted reward is zero and the dopaminergic neurons firing corresponding to a positive error; the reward exceeded expectation. The bottom panel shows what happens after conditioning if the reward is not received, since a reward is predicted this is a negative error and depresses firing. The middle panel shows that the consequence of conditioning is to advance the dopamine firing forward in time. [Figure taken from [5]].

References




Figure 6: In **A1** the change in strength of EPSPs is shown plotted against time, the synapse with closed circles has been stimulated to induce long term potentiation, the one with open circles is a control. In **A2** the same thing is done, but with dopamine receptors blocked during the stimulation. [Picture from [9] which in turn adapted it from [10]]

The VTA and hippocampus

Back in Fig. 3 we saw that VTA also projects to the hippocampus; we have previously examined memory in the hippocampus so it is interesting to speculate what the role of this projection is. In [9] it is argued that there is a signal from VTA to hippocampal in which dopamine is used to mark novel and salient stimuli, prompting their encoding in memory. In this model a new stimulus is tested in CA1 or PRC to see if it is already stored in hippocampus, if it isn't, this is communicated to VTA where the local activity will test how novel, salient or surprising this stimulus is, this is a role akin to, but slightly different from the role described above in classical conditioning. If it is novel, activity in the dopaminergic neurons increased long term plastic changes in hippocampus. Evidence for this is found in [10] where it is shown that blocking dopamine receptors in hippocampus reduces learning.

References

- [1] Rescorla RA and Wagner AR. (1972) A theory of Pavlovian conditioning: Variations in the effectiveness of reinforcement and nonreinforcement, *Classical Conditioning II*, Black, AH and Prokasy WF, Eds., 64–99. Appleton-Century-Crofts.
- [2] Dayan P and Abbott FA. (2001) *Theoretical neuroscience*. MIT press (Cambridge, MA).
- [3] Miller RR, Barnet RC and Grahame NJ. (1995) Assessment of the Rescorla-Wagner model. *Psychological bulletin* 117: 363–86
- [4] Azorlosa JL and Cicala GA. (1986) Blocking of conditioned suppression with 1 or 10 compound trials. *Animal Learning and Behavior* 14: 163–7.
- [5] Schultz W, Dayan P and Montague PR. (1997) A neural substrate of prediction and reward. *Science* 275: 1593–9.
- [6] Sutton R (1988) Learning to predict by the methods of temporal differences. *Machine Learning* 3: 9–44.

References

- [7] Sutton R and Barto S (1998). Reinforcement Learning. MIT Press (Cambridge, MA).
- [8] Barto AG. (1995) Adaptive critics and the basal ganglia. Models of information processing in the basal ganglia. Eds. Houk JC, Davis JL and Beiser DG. MIT press (Cambridge, MA).
- [9] Lisman JE and Grace AA (2005) The hippocampal-VTA loop: controlling the entry of information into long-term memory. *Neuron* 46: 703–13.
- [10] Morris RG, Moser EI, Riedel G, Martin SJ, Sandin J, Day M and O’Carroll C (2003) Elements of a neurobiological theory of the hippocampus: the role of activity-dependent synaptic plasticity in memory. *Philosophical Transactions of the Royal Society*. 358: 773–786.